



Machine Learning with Graphs (MLG)

# RecSys: Bayesian Personalized Ranking (BPR) Loss

Learning to rank items for users

Cheng-Te Li (李政德)

Institute of Data Science

National Cheng Kung University


chengte@mail.ncku.edu.tw




# Explicit vs. Implicit Feedback

- Learning from explicit feedback (e.g., ratings)
  - CF, MF, and FM
  - → However, at most time, users provide no ratings
- Much easier to collect implicit feedback
  - Clicks on pages and URLs
  - Purchases
  - View times
  - → Already available in log files at the web servers
- Can we learn personalized ranking from implicit data for recommendation?

# Ranking on Items




# Ranking on Items




# Personalized Ranking

- Goal: provide a user with a ranked list of items
  - Ranking is inferred from implicit behaviors



# Filling with “0” ?

- Implicit feedback contains only positive classes, the remaining is a mixture of unknown and negative
  - Filling 0 for unknown/negative, then do CF/MF/FM?
  - This approach tend to predict 0, it cannot work!



# Ranking Settings

- Only know the ranking between two items but not exactly scores
- Reformat each user-item pairs
  - Training contains both positive and negative
  - Missing to be ranked
- Training data

$$D_S: U \times I \times I$$

$$D_S = \{(u, i, j) \mid i \in I_u^+ \wedge j \in I \setminus I_u^+\}$$

$D_S$ : set of triples that user  $u$  likes item  $i$  more than item  $j$

Implicit feedback  
(only positive)

	$i_1$	$i_2$	$i_3$	$i_4$	
$u_1$	?	+	+	?	user ↑ ↓
$u_2$	+	?	?	+	
$u_3$	+	+	?	?	
$u_4$	?	?	+	+	
$u_5$	?	?	+	?	
	← item →				

$$S: U \times I$$

$$I_u^+ := \{i \in I : (u, i) \in S\}$$

$$U_i^+ := \{u \in U : (u, i) \in S\}$$

$u_1: i >_{u_1} j$

	$i_1$	$i_2$	$i_3$	$i_4$	
$j_1$		+	+	?	item ↑ ↓
$j_2$	-		?	-	
$j_3$	-	?		-	
$j_4$	?	+	+		
	← item →				


...

$u_5: i >_{u_5} j$

	$i_1$	$i_2$	$i_3$	$i_4$	
$j_1$		?	+	?	item ↑ ↓
$j_2$	?		+	?	
$j_3$	-	-		-	
$j_4$	?	?	+		
	← item →				

# AUC behind BPR

- Goal: rank item  $i$  higher than item  $j$  (i.e.,  $i >_u j$ )
  - $\rightarrow$  Given item  $i$  higher score than  $j$
- AUC = Area under ROC curve
  - $0 \leq \text{AUC} \leq 1$
- Correctly ranking  $i >_u j$   
 $\approx$  Maximize AUC
- When  $\text{AUC} > 0.5$ :  
the prediction is above “slope=1”  
 $\rightarrow$  Tend to predict  $i >_u j$  and  
better than random guess





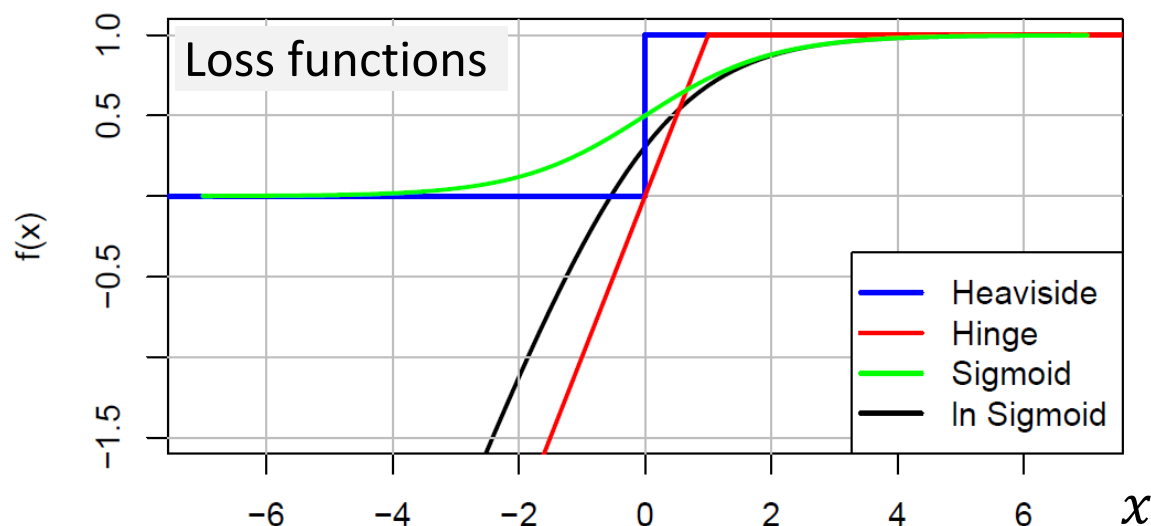
# Optimize AUC

$$\text{AUC}(u) := \frac{1}{|I_u^+| |I \setminus I_u^+|} \sum_{i \in I_u^+} \sum_{j \in |I \setminus I_u^+|} \delta(\hat{x}_{uij} > 0)$$

$\hat{x}_{uij}$ : any real-valued function that gives the ranking order between  $i$  and  $j$

$$\text{AUC} := \frac{1}{|U|} \sum_{u \in U} \text{AUC}(u) \quad \delta(x > 0) = H(x) := \begin{cases} 1, & x > 0 \\ 0, & \text{else} \end{cases}$$

Heaviside function (單位階梯函數)



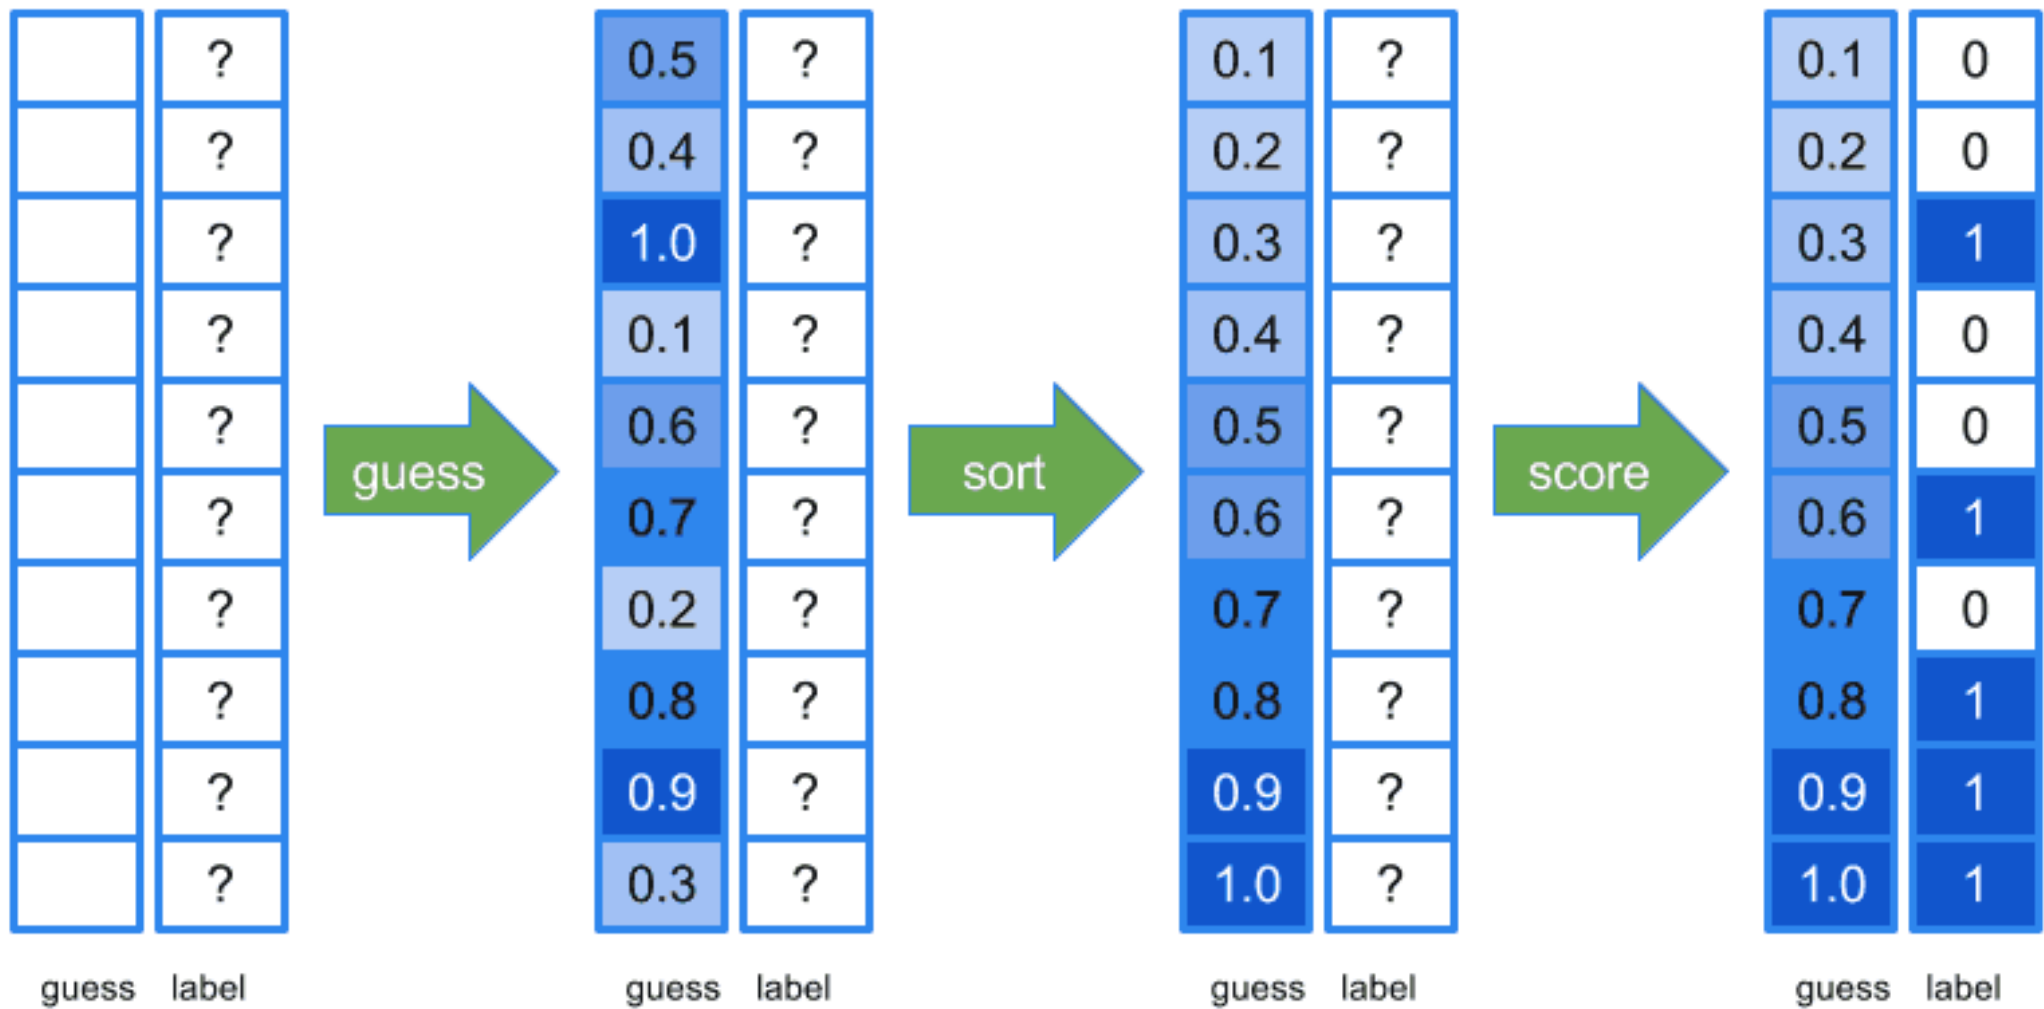
→ non-differentiable (不可微分)

→ ROC is NOT a differentiable smooth curve


i.e., ROC curve is stepwise shape

→ We cannot use gradient descent to optimize AUC score


# AUC Explained




# AUC Explained




# AUC Explained




# AUC Explained




# AUC Explained



# AUC Explained




# AUC Explained





# AUC Explained



# AUC Explained

0.1	0
0.2	0
0.3	1
0.4	0
0.5	0
0.6	1
0.7	0
0.8	1
0.9	1
1.0	1

guess label




threshold


Threshold: 0.1

TPR: 5 / 5


FPR: 4 / 5




# AUC Explained



# AUC Explained



# AUC Explained



# Bayesian Personalized Ranking

- Goal: optimize the posterior probability

$$p(\Theta | >_u) \propto p(>_u | \Theta)p(\Theta)$$

$\Theta$ : parameters of any model (e.g., MF)

$>_u$ : the desired ranking of items for user  $u$

- Assume every user-item pair  $(u, i)$  is independent of each other, and users' preference are also independent

$$\prod_{u \in U} p(>_u | \Theta) = \prod_{(u, i, j) \in D_S} p(i >_u j | \Theta)$$


# Bayesian Personalized Ranking

- Use **sigmoid** function to estimate probability  $p(>_u | \Theta)$ 
  - For gradient descent, the objective needs to be differentiable

Rather than  
Heaviside function

$$p(i >_u j | \Theta) := \sigma(\hat{x}_{uij}(\Theta))$$

$$\sigma(x) := \frac{1}{1 + e^{-x}}$$



For regularization and for taking log:

# BPR Optimization

$\therefore$  assume:  $\Theta \sim N(0, \lambda_{\Theta} I)$

$$\begin{aligned} \text{BPR-OPT} &:= \ln p(\Theta | >_u) \\ &= \ln p(>_u | \Theta) p(\Theta) \\ &= \ln \prod_{(u,i,j) \in D_S} \sigma(\hat{x}_{uij}) p(\Theta) \\ &= \sum_{(u,i,j) \in D_S} \ln \sigma(\hat{x}_{uij}) + \ln p(\Theta) \\ &= \sum_{(u,i,j) \in D_S} \ln \sigma(\hat{x}_{uij}) - \lambda_{\Theta} ||\Theta||^2 \end{aligned}$$
$$p(\Theta) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{\|\Theta\|^2}{2\sigma^2}\right)$$
$$\ln p(\Theta) = -\lambda \|\Theta\|^2$$



# BPR Optimization

- Gradient descent 
$$\frac{\partial \text{BPR-Opt}}{\partial \Theta} = \sum_{(u,i,j) \in D_S} \frac{\partial}{\partial \Theta} \ln \sigma(\hat{x}_{uij}) - \lambda_{\Theta} \frac{\partial}{\partial \Theta} \|\Theta\|^2$$
$$\propto \sum_{(u,i,j) \in D_S} \frac{-e^{-\hat{x}_{uij}}}{1 + e^{-\hat{x}_{uij}}} \cdot \frac{\partial}{\partial \Theta} \hat{x}_{uij} - \lambda_{\Theta} \Theta$$

```
1: procedure LEARNBPR( $D_S, \Theta$ )
2:   initialize  $\Theta$ 
3:   repeat
4:     draw  $(u, i, j)$  from  $D_S$  Bootstrap sampling
5:      $\Theta \leftarrow \Theta + \alpha \left( \frac{e^{-\hat{x}_{uij}}}{1 + e^{-\hat{x}_{uij}}} \cdot \frac{\partial}{\partial \Theta} \hat{x}_{uij} + \lambda_{\Theta} \cdot \Theta \right)$ 
6:   until convergence
7:   return  $\hat{\Theta}$ 
8: end procedure
```

# Matrix Factorization with BPR (BPR-MF)

- BPR is an ranking-based optimization idea, NOT RecSys
  - Can be adopted for existing RecSys models and deep models that can produce real value  $\hat{x}_{ui}$  for a user-item pair  $(u, i)$
- MF:  $\hat{X} = WH^T$ , i.e.,  $\hat{x}_{ui} = \sum_{f=1}^k w_{uf} h_{if}$ 
  - Model parameters  $\Theta = \{W, H\}$  (latent features of users / items)

$$BPR = \sum_{(u,i,j) \in D_S} \ln \sigma(\hat{x}_{uij}) - \lambda_{\Theta} \|\Theta\|^2$$

$\sigma$ : sigmoid

$\hat{x}_{uij}$ : any real-valued function that gives the ranking order between  $i$  and  $j$

- Define the predicted ranking order  $\hat{x}_{uij} = \hat{x}_{ui} - \hat{x}_{uj}$

$$\text{AUC}(u) := \frac{1}{|I_u^+| |I \setminus I_u^+|} \sum_{i \in I_u^+} \sum_{j \in I \setminus I_u^+} \delta(\hat{x}_{uij} > 0)$$

NOT to **regress** a single value  $\hat{x}_{ui}$ , but to **classify**  $\hat{x}_{ui} - \hat{x}_{uj}$

$$\text{BPR-MF}^{BPR} = \sum_{(u,i,j) \in D_S} \ln \sigma(\hat{x}_{uij}) - \lambda_{\Theta} \|\Theta\|^2$$

- Training data:  $D_S = \{(u, i, j) \mid i \in I_u^+ \wedge j \in I \setminus I_u^+\}$
- The predicted ranking order  $\hat{x}_{uij} = \hat{x}_{ui} - \hat{x}_{uj}$

$$\hat{x}_{ui} = \sum_{f=1}^k w_{uf} h_{if}$$

$$\sigma(x) = \frac{1}{1 + \exp(-x)}$$

```

1  # compute x_uij
2  Wu = W[u].view(1, W[u].size()[0])
3  x_uij = torch.mv(Wu, H[i]) - torch.mv(Wu, H[j])
4
5  # compute e^-x_uij / (1 + e^-x_uij)
6  exp_x = np.exp(-x_uij)
7  partial_BPR = exp_x / (1 + exp_x)
8
9  # 對第 1 ~ k 個 feature 更新
10 for f in range(k):
11     W[u][f] -= lr * (partial_BPR * (H[i][f] - H[j][f]) + rr * W[u][f])
12     H[i][f] -= lr * (partial_BPR * W[u][f] + rr * H[i][f])
13     H[j][f] -= lr * (partial_BPR * (-W[u][f]) + rr * H[i][f])


```

$$\frac{\partial}{\partial \theta} \hat{x}_{uij} = \begin{cases} (h_{if} - h_{jf}) & \text{if } \theta = w_{uf}, \\ w_{uf} & \text{if } \theta = h_{if}, \\ -w_{uf} & \text{if } \theta = h_{jf}, \\ 0 & \text{else} \end{cases}$$


[https://github.com/leafinity/gradient\\_dscent\\_svd/blob/master/bpr.ipynb](https://github.com/leafinity/gradient_dscent_svd/blob/master/bpr.ipynb)

# BPR Performance

Online shopping: Rossmann



Video Rental: Netflix



# Short Summary

- BPR is ranking-aware optimization
  - Separate unknown from missing in “0”s
  - Use gradient descent to optimize non-differential AUC
  - Bootstrap sampling for  $D_S = \{(u, i, j)\}$
  - Can be combined with different models
- Better selection of  $D_S$ 
  - E.g., session-based selection
- We will see BPR many times soon



U1,A,B  
U1,A,D  
U1,A,E  
U1,C,B  
U1,C,D  
U1,C,E  
U1,F,H  
U1,F,I  
U1,G,H  
U1,G,I  
U1,J,H  
U1,J,I

# References

- S. Rendle et al. “BPR: Bayesian Personalized Ranking from Implicit Feedback” UAI 2009 2480 cites
- <https://medium.com/@radleaf/bpr-and-recommendation-system-3d9a3975c132>
- <https://blog.csdn.net/sigmaeta/article/details/80517828>
- <https://www.cnblogs.com/wkang/p/10217172.html>
- <https://www.biaodianfu.com/bpr.html>

# BPR Packages/Codes

- Case Recommender
  - <https://github.com/caserec/CaseRecommender>
- **Spotlight** [complete, recommended!!]
  - <https://github.com/maciejkula/spotlight>
- **LightFM** [recommended!!]
  - <https://github.com/lyst/lightfm/>
- RecSys tutorial
  - [https://github.com/MaurizioFD/RecSys\\_Course\\_AT\\_PoliMi](https://github.com/MaurizioFD/RecSys_Course_AT_PoliMi)
- **NeuRec** [complete, recommended!!]
  - <https://github.com/wubinzzu/NeuRec>